

# 8D\_Solution

April 10, 2022

## 0.1 Task-D: Collinear features and their effect on linear models

```
[1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import SGDClassifier
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
[2]: data = pd.read_csv('task_d.csv')
```

```
[3]: data.head()
```

```
[3]:
```

	x	y	z	x*x	2*y	2*z+3*x*x	w \
0	-0.581066	0.841837	-1.012978	-0.604025	0.841837	-0.665927	-0.536277
1	-0.894309	-0.207835	-1.012978	-0.883052	-0.207835	-0.917054	-0.522364
2	-1.207552	0.212034	-1.082312	-1.150918	0.212034	-1.166507	0.205738
3	-1.364174	0.002099	-0.943643	-1.280666	0.002099	-1.266540	-0.665720
4	-0.737687	1.051772	-1.012978	-0.744934	1.051772	-0.792746	-0.735054

	target
0	0
1	0
2	0
3	0
4	0

```
[4]: X = data.drop(['target'], axis=1).values
Y = data['target'].values
```

## 0.2 Task 1 : Logistic Regression

## 1. Finding the Correlation between the features

```
[5]: corr = data[data.columns[:-1]].corr()  
corr
```

```
[5]:
```


	x	y	z	x*x	2*y	2*z+3*x*x	\
x	1.000000	-0.205926	0.812458	0.997947	-0.205926	0.996252	
y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	
z	0.812458	-0.602663	1.000000	0.807137	-0.602663	0.847163	
x*x	0.997947	-0.209289	0.807137	1.000000	-0.209289	0.997457	
2*y	-0.205926	1.000000	-0.602663	-0.209289	1.000000	-0.261123	
2*z+3*x*x	0.996252	-0.261123	0.847163	0.997457	-0.261123	1.000000	
w	0.583277	-0.401790	0.674486	0.583803	-0.401790	0.606860	

	w
x	0.583277
y	-0.401790
z	0.674486
x*x	0.583803
2*y	-0.401790
2*z+3*x*x	0.606860
w	1.000000

```
[6]: plt.figure(figsize=(8,6))  
plt.title("Correlation between the features")  
sns.heatmap(corr,cmap="BrBG",annot = True)
```

```
[6]: <AxesSubplot:title={'center':'Correlation between the features'}>
```



## 2. Finding the best model for the given data

```
[7]: alpha = np.logspace(-5,8,10)
      print("Alpha values:",alpha)
      param_grid={'C':alpha}
```

```
Alpha values: [1.00000000e-05 2.78255940e-04 7.74263683e-03 2.15443469e-01
 5.99484250e+00 1.66810054e+02 4.64158883e+03 1.29154967e+05
 3.59381366e+06 1.00000000e+08]
```

```
[8]: LR = LogisticRegression(max_iter=10000, tol=0.1, random_state = 0)
      LR = GridSearchCV(LR,param_grid,cv=5)
      LR.fit(X,Y)
      print("Best alpha is :",LR.best_params_['C'])
```

```
Best alpha is : 1e-05
```

## 3. Getting the weights with the original data

```
[9]: best_model = LogisticRegression(max_iter=10000, tol=0.1,C=1e-05,random_state = 0)
      best_model.fit(X,Y)
```

```
acc = best_model.score(X,Y)
print("Accuracy of the model is {}".format(acc*100))
```

Accuracy of the model is 100.0%

```
[10]: weight = best_model.coef_[0]
print("Weights =",weight)
```

```
Weights = [ 0.0003637 -0.00034492  0.00048441  0.00035935 -0.00034492
0.0003819
 0.00032049]
```

#### 4. Modifying original data

```
[11]: X_New = np.random.normal(0,10**-2,X.shape)
best_model1 = LogisticRegression(max_iter=10000, tol=0.1,C=1e-05,random_state =
↳0)
best_model1.fit(X_New,Y)

acc1 = best_model1.score(X_New,Y)
print("Accuracy of the model with noise is {}".format(acc1*100))
```

Accuracy of the model with noise is 62.0%

```
[12]: weight1 = best_model1.coef_
print("Weights =",weight1)
```

```
Weights = [[-1.03721910e-08  1.24072634e-06 -1.97814951e-07 -1.38853930e-07
-5.89260746e-07 -6.96008720e-07  1.37269248e-07]]
```

#### 5. Checking deviations in metric and weights

```
[13]: acc_diff = abs(acc - acc1)
print("Difference in Accuracy : {}".format(np.round(acc_diff*100)))

weight_diff = abs(weight-weight1)
print("Difference in weights:",weight_diff)

weight_diff_per = weight_diff * 100
print("percentage Difference in weights:",weight_diff_per)
```

```
Difference in Accuracy : 38.0%
Difference in weights: [[0.00036371 0.00034616 0.00048461 0.00035949 0.00034433
0.0003826
 0.00032035]]
percentage Difference in weights: [[0.03637143 0.03461645 0.04846054 0.0359488
0.03443345 0.03825971
 0.03203492]]
```

```
[14]: top4 = np.argsort(-weight_diff_per)[:,:4]
      features = data.columns
      print("Top 4 features are :",features[top4])
```

Top 4 features are : [['z' '2\*z+3\*x\*x' 'x' 'x\*x']]

### 0.3 Task2 : Linear SVM

#### 1.Finding the best SVM model for the given data

```
[17]: alpha_svm = np.logspace(-5, 8, 10)
      print("Alphas = ",alpha_svm)
      param_grid={'C':alpha_svm}
      svm = SVC(kernel="linear",random_state = 0)
```

Alphas = [1.00000000e-05 2.78255940e-04 7.74263683e-03 2.15443469e-01  
5.99484250e+00 1.66810054e+02 4.64158883e+03 1.29154967e+05  
3.59381366e+06 1.00000000e+08]

```
[16]: model_svm = GridSearchCV(svm, param_grid, cv=5)
      model_svm.fit(X,Y)
      print("Best alpha is:",model_svm.best_params_['C'])
```

Best alpha is: 0.007742636826811269

#### 2.Getting the weights with the original data

```
[18]: best_model_svm=SVC(kernel='linear',C=0.007742636826811269,random_state=0)
      best_model_svm.fit(X,Y)

      acc_svm = best_model_svm.score(X,Y)
      print("Accuracy of the model is {}".format(acc_svm*100))
```

Accuracy of the model is 100.0%

```
[19]: weight_svm = best_model_svm.coef_
      print("Weights =",weight_svm)
```

Weights = [[ 0.16056222 -0.20788705 0.32826166 0.14998082 -0.20788705  
0.17461587  
0.13401176]]

#### 3.Modifying original data

```
[20]: X_New_svm = X + np.random.normal(0,10**-2,X.shape)
      best_model_svm = SVC(kernel='linear',C=0.007742636826811269,random_state=0)
      best_model_svm.fit(X_New_svm,Y)

      acc1_svm = best_model_svm.score(X_New_svm,Y)
      print("Accuracy of the model is {}".format(acc1_svm*100))
```

Accuracy of the model is 100.0%

```
[21]: weight1_svm = best_model_svm.coef_  
print("Weights =",weight1_svm)
```

```
Weights = [[ 0.16112764 -0.20612367  0.32858142  0.15085447 -0.20595574  
0.17495079  
0.13471504]]
```

#### 4. Checking deviations in metric and weights

```
[22]: acc_diff_svm = abs(acc1_svm - acc_svm)  
print("Difference in Accuracy : {}".format(np.round(acc_diff_svm*100)))  
  
weight_diff_svm = abs(weight1_svm - weight_svm)  
print("Difference in weights:",weight_diff_svm)  
  
weight_diff_per_svm = weight_diff_svm * 100  
print("Percentage difference in weights",weight_diff_per_svm)  
  
top4_svm = np.argsort(-weight_diff_per_svm)[:,:4]  
features = data.columns  
print("Top 4 features are:",features[top4_svm])
```

```
Difference in Accuracy : 0.0%  
Difference in weights: [[0.00056542 0.00176338 0.00031976 0.00087365 0.00193131  
0.00033492  
0.00070329]]  
Percentage difference in weights [[0.05654215 0.17633837 0.03197599 0.08736534  
0.19313099 0.03349245  
0.07032868]]  
Top 4 features are: [['2*y' 'y' 'x*x' 'w']]
```

#### 0.4 Observations :

- Multi-collinearity doesnot affect both Logistic regression and Linear SVM.
- In case of Logistic Regression when we add the noise, the accuracy of the model is decreased.
- Noise does not affect the accuracy of Linear SVM model.